

# Interpretable Machine Learning, Part 4: SHAP

# Shapely Values

Shapely values are a concept developed in 1951 in the field of game theory.

**Big Idea:** If you have a coalition of players all working together, and they achieve a certain payout, how do we divide that payout up into the individual contributions of each player?

That is, how important is each individual player?

# Shapely Values

What is the connection to predictive models?

Let's say we are trying to predict the price of a house based on *sqft\_living*, *grade*, *waterfront*, and *bedrooms*, and the average home price in our dataset is \$540,000.

We have a 3 bedroom, 3,560 square foot home, with grade 8, with waterfront 2, for which to model predicts a price of \$768,000.

Here the “payout” is the difference between our prediction and the average price ( $\$768,000 - \$540,000 = \$228,000$ )

We need to determine how to divide the \$228,000 difference among our 4 features.


# Shapely Values - Mathematical Definition

$$\varphi_i(v) = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|! (n - |S| - 1)!}{n!} (v(S \cup \{i\}) - v(S))$$


Shapely  
value for  
feature  $i$




Sum over all  
subsets of  
features that  
don't include  
feature  $i$




Weight for that  
subset of  
features (large  
and small set of  
features get  
more weight)



Predicted  
value with  
feature  $i$



Predicted  
value without  
feature  $i$



# Shapely Values

The formula from the previous slide involved calculating the prediction with and without certain features.

But, we cannot actually remove the feature, since the model expects us to pass in a value for it.

To “remove” the feature, we take the same approach as permutation importance: substitute in a random value instead.

If we repeat this a large number of times and take the average, we will have a pretty good approximation of the true Shapely value.

# Shapely Values

| Feature     | Including Bedrooms     | Bedrooms "Removed" |
|-------------|------------------------|--------------------|
| sqft_living | 3560                   | 3560               |
| grade       | 8                      | 8                  |
| waterfront  | 0                      | 0                  |
| bedrooms    | 3                      | 2                  |
| PREDICTION: | <b>\$768,000</b>       | <b>\$787,000</b>   |
|             | MARGINAL CONTRIBUTION: | <b>-\$19,000</b>   |

First, we substitute in random values for the number of bedrooms in order to "remove" that feature.

# Shapely Values

| Feature     | Including Bedrooms     | Bedrooms "Removed" |
|-------------|------------------------|--------------------|
| sqft_living | 3560                   | 3560               |
| grade       | 8                      | 8                  |
| waterfront  | 0                      | 0                  |
| bedrooms    | 3                      | 4                  |
| PREDICTION: | <b>\$768,000</b>       | <b>\$1,045,00</b>  |
|             | MARGINAL CONTRIBUTION: | <b>-\$277,000</b>  |

# Shapely Values

| Feature     | Including Bedrooms     | Bedrooms "Removed" |
|-------------|------------------------|--------------------|
| sqft_living | 2480                   | 2480               |
| grade       | 8                      | 8                  |
| waterfront  | 0                      | 0                  |
| bedrooms    | 3                      | 2                  |
| PREDICTION: | <b>\$547,000</b>       | <b>\$727,000</b>   |
|             | MARGINAL CONTRIBUTION: | <b>-\$180,000</b>  |

Here, we have  
"removed"  
the *sqft\_living*  
feature



# Shapely Values

| Feature     | Including Bedrooms     | Bedrooms "Removed" |
|-------------|------------------------|--------------------|
| sqft_living | 2480                   | 2480               |
| grade       | 8                      | 8                  |
| waterfront  | 0                      | 0                  |
| bedrooms    | 3                      | 4                  |
| PREDICTION: | <b>\$547,000</b>       | <b>\$557,000</b>   |
|             | MARGINAL CONTRIBUTION: | <b>-\$10,000</b>   |

# Shapely Values

| Feature     | Including Bedrooms     | Bedrooms "Removed" |
|-------------|------------------------|--------------------|
| sqft_living | 1500                   | 1500               |
| grade       | 9                      | 9                  |
| waterfront  | 0                      | 0                  |
| bedrooms    | 3                      | 4                  |
| PREDICTION: | <b>\$671,000</b>       | <b>\$772,000</b>   |
|             | MARGINAL CONTRIBUTION: | <b>-\$101,000</b>  |


Here, we have  
"removed"  
the *sqft\_living*  
and the *grade*  
features

# Shapely Values

Average home price: 540088.14

Predicted home price: 768655.22

Difference from Average: 228567.08




Repeating this a large number of times and for each feature, we get this result.

# Shapely Values

Average home price: 540088.14

Predicted home price: 768655.22

Difference from Average: 228567.08




The *bedrooms* feature decreased the prediction value by about \$40,000

# Shapely Values

Average home price: 540088.14

Predicted home price: 768655.22


Difference from Average: 228567.08



The *sqft\_living* feature increased the prediction value by about \$270,000

# Shapely Values

We can also see the results as a *force plot*, which takes the previous plot and displays it in a stacked format.



# SHAP

In Python, we'll be using the *shap* library, which uses a method called SHAP (**SH**apely **A**dditive ex**P**lanations).

For each prediction, SHAP builds an “explanations model”  $g$ , where  $\phi_j$  is the feature attribution for feature  $j$ , and  $z_j'$  indicates whether that feature is present or absent.

$$g(z') = \phi_0 + \sum_{j=1}^M \phi_j z_j'$$

# SHAP

For more details, see the original paper:

<https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions.pdf>




# SHAP

There are two “flavors” of SHAP available in the *shap* library:

- KernelSHAP
  - This is essentially the procedure described above where we randomly replace certain features and see the effect on the predicted values
  - Computes an *estimate* of the Shapely values for each feature
- TreeSHAP
  - A variant of SHAP for tree-based models (random forests, and gradient-boosted trees, like XGBoost)
  - Computes *exact* Shapely values for each feature

# SHAP Values


Average home price: 540088.14  
Predicted home price: 1013793.76  
Difference from Average: 473705.62



Another example.  
This time, a home  
with predicted  
price of  
\$1,014,000.  
Where did the  
extra \$470,000  
come from?

# SHAP Values

Average home price: 540088.14  
Predicted home price: 1013793.76  
Difference from Average: 473705.62




A great deal of it, almost \$350,000, came from the fact that the grade was so high.

# SHAP Values

Average home price: 540088.14

Predicted home price: 1013793.76


Difference from Average: 473705.62



The living space  
added about  
\$100,000

# SHAP Values


Average home price: 540088.14  
Predicted home price: 1013793.76  
Difference from Average: 473705.62



Since it was not a waterfront property, the predicted price dropped by about \$5,000.


# SHAP Values

We can also see the results as a *force plot*, which takes the previous plot and displays it in a stacked format.




# SHAP Values

We can also see how different features impact predictions across the entire dataset, similarly to a partial dependence plot.



# SHAP Values

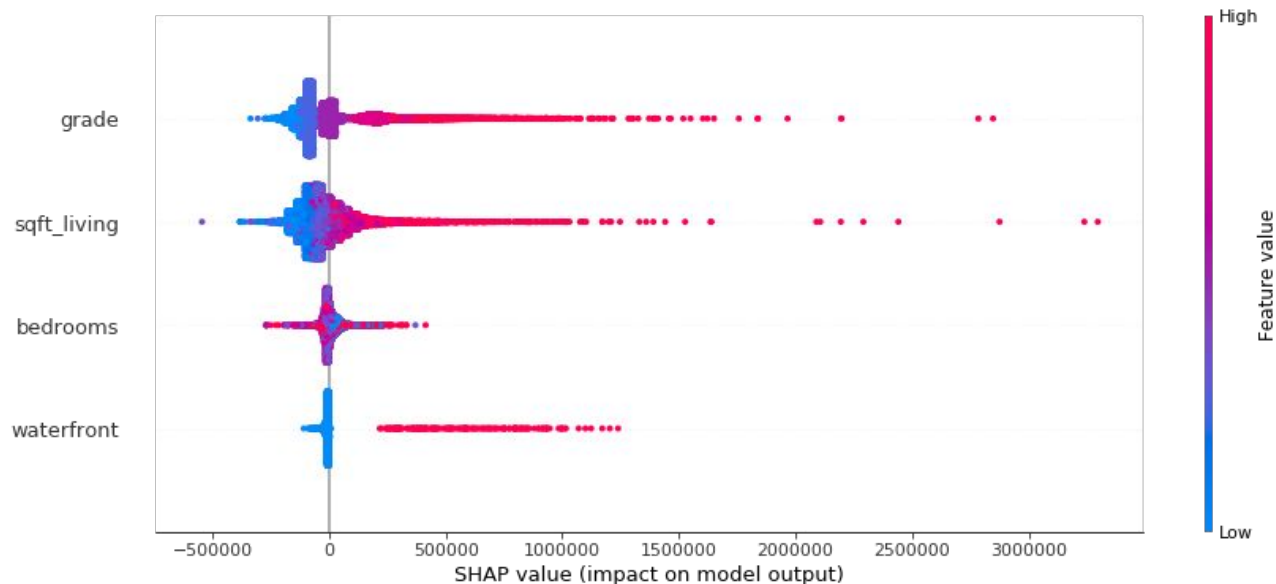
We can also see how different features impact predictions across the entire dataset, similarly to a partial dependence plot.





# SHAP Values

Or see the effect of each feature across the entire dataset.



# SHAP Values

Or see the effect of each feature across the entire dataset.

